

軟體分析與最佳化 期末 Stage 1

報告人：

612410017 林靖紳

612410066 蔡宏遠

Outline

- Program Introduction
- Workload Structure
- Evaluation Environment
- Compile and Run
- Experimental Result
- Reference

Program Introduction

Parsec - Dedup

- **Objective:**

- Achieve data compression by detecting and eliminating duplicate blocks within input data.

- **Process:**

- **Input Data Chunking:**

- Divide input data into fixed-sized blocks, serving as the fundamental processing units

- **Duplicate Block Detection:**

- Utilize hash functions or other digest algorithms to compare content between different blocks, identifying duplicate blocks.

- **Deduplication Process:**

- Retain only one copy of duplicate blocks, referencing it as needed, thereby significantly reducing storage requirements.

- **Output Compression:**

- The output data from Dedup is compressed, containing only unique blocks, resulting in decreased storage needs and improved efficiency during data transmission.

Parsec - Dedup Performance Measurement

- **test**: Minimal input to verify that programs are executable
- **simdev**: Very small input which causes code execution comparable to a typical input for this program. Intended for microarchitectural simulator development
- **simsmall**: Small input for performance measurements with microarchitectural simulators
- **simmedium**: Medium-sized input for performance measurements with microarchitectural simulators
- **simlarge**: Large-sized input for performance measurements with microarchitectural simulators
- **native**: Very large input intended for large-scale experiments on real machines

Workload structure

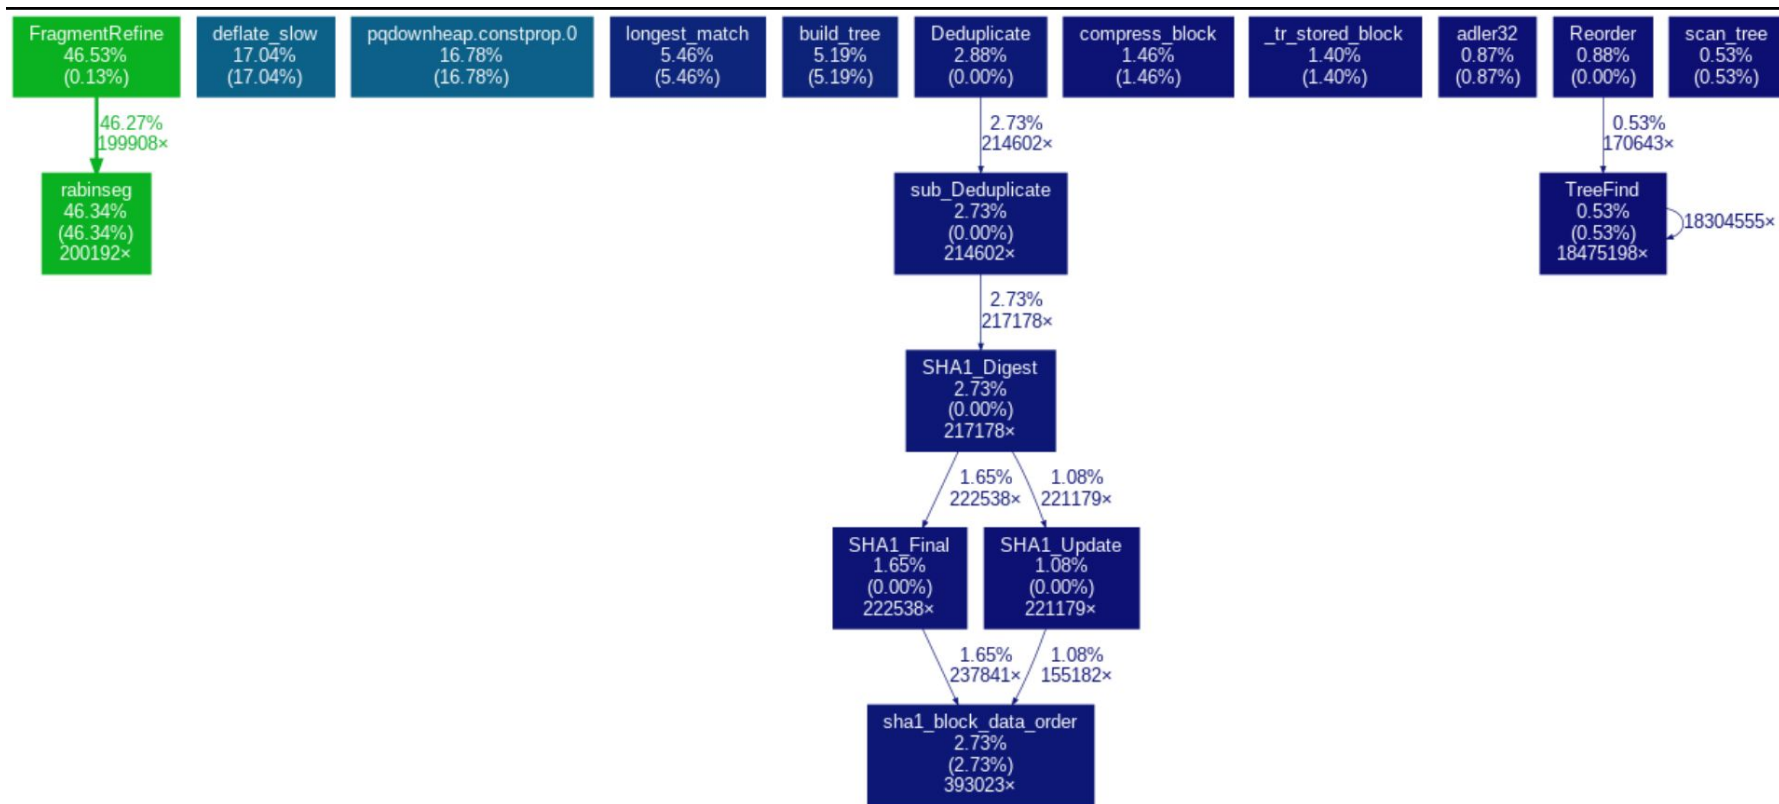
gprof 執行

```
ashen@Stephanie-Lin:~/parsec-benchmark-master/pkgs/kernels/dedup/inst/amd64-linux.gcc/bin$ gprof ./dedup
Flat profile:
```

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
46.37	6.96	6.96	200192	34.77	34.77	rabinseg
17.06	9.52	2.56				deflate_slow
16.79	12.04	2.52				pqdownheap.constprop.0
5.46	12.86	0.82				longest_match
5.20	13.64	0.78				build_tree
2.73	14.05	0.41	393023	1.04	1.04	sha1_block_data_order
1.47	14.27	0.22				compress_block
1.40	14.48	0.21				_tr_stored_block
0.87	14.61	0.13				adler32
0.53	14.69	0.08	170643	0.47	0.47	TreeFind
0.53	14.77	0.08				scan_tree
0.40	14.83	0.06				fill_window
0.17	14.86	0.03	2816537	0.01	0.01	ringbuffer_isEmpty
0.13	14.88	0.02	200853	0.10	0.11	DeleteMin
0.13	14.89	0.02				FragmentRefine
0.13	14.91	0.02				_tr_flush_block
0.13	14.94	0.02				deflateReset
0.13	14.96	0.02				read_tree

gprof2dot



Evaluation Environment

CPU Information

```
ashen@Stephanie-Lin:~$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Address sizes:          39 bits physical, 48 bits virtual
Byte Order:             Little Endian
CPU(s):                 12
On-line CPU(s) list:    0-11
Vendor ID:              GenuineIntel
Model name:             11th Gen Intel(R) Core(TM) i5-11500 @ 2.70GHz
CPU family:             6
Model:                  167
Thread(s) per core:     2
Core(s) per socket:     6
Socket(s):              1
Stepping:               1
CPU max MHz:            4600.0000
CPU min MHz:            800.0000
BogoMIPS:               5424.00
```

Memory

```
ashen@Stephanie-Lin:~$ free -h
```

	total	used	free	shared	buff/cache	available
Mem:	31Gi	2.7Gi	3.9Gi	1.1Gi	24Gi	26Gi
Swap:	2.0Gi	0B	2.0Gi			

```
ashen@Stephanie-Lin:~$
```

OS Version

```
ashen@Stephanie-Lin:~$ lsb_release -a
```

No LSB modules are available.

Distributor ID: Ubuntu

Description: Ubuntu 22.04.3 LTS

Release: 22.04

Codename: jammy

gcc version

```
ashen@Stephanie-Lin:~$ gcc --version
gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
Copyright (C) 2021 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

icc version

```
ashen@Stephanie-Lin:~$ icc --version
icc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is dep
lease in the second half of 2023. The Intel(R) oneAPI DPC++/C++ Co
ving forward. Please transition to use this compiler. Use '-diag-d
icc (ICC) 2021.10.0 20230609
Copyright (C) 1985-2023 Intel Corporation. All rights reserved.
```

Compile and Run

Dedup configure

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ ./configure
[sudo] password for ashen:
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
build-essential is already the newest version (12.9ubuntu3).
libglu1-mesa-dev is already the newest version (9.0.2-1).
libxi-dev is already the newest version (2:1.8-1build1).
libxmu-dev is already the newest version (2:1.1.3-3).
m4 is already the newest version (1.4.18-5ubuntu2).
x11proto-xext-dev is already the newest version (2021.5-1).
libtbb-dev is already the newest version (2021.5.0-7ubuntu2).
0 upgraded, 0 newly installed, 0 to remove and 8 not upgraded.
+ sim=true
+ verbose=
+ getopts Sv OPT
+ shift 0
+ download_dir=.
+ outdir=.
+ url_base=http://parsec.cs.princeton.edu/download/3.0
+ basenames=
+ basenames=' parsec-3.0-core.tar.gz'
+ true
+ basenames=' parsec-3.0-core.tar.gz parsec-3.0-input-sim.tar.gz'
+ mkdir -p .
+ for basename in $basenames
+ '[' '!' -f ./parsec-3.0-core.tar.gz ']'
+ wget -nc -P . http://parsec.cs.princeton.edu/download/3.0/parsec-3.0-core.tar.gz
--2023-11-16 14:51:51-- http://parsec.cs.princeton.edu/download/3.0/parsec-3.0-core.tar.gz
Resolving parsec.cs.princeton.edu (parsec.cs.princeton.edu)... failed: Name or service not known.
wget: unable to resolve host address 'parsec.cs.princeton.edu'
ashen@Stephanie-Lin:~/parsec-benchmark-master$ source env.sh
```


Dedup build with gcc -O3

```
73 # Arguments to use
74 export CFLAGS="-DUNIX -O3 -funroll-loops -fprefetch-loop-arrays ${PORTABILITY_FLAGS} -pg"
75 export CXXFLAGS="-DUNIX -O3 -funroll-loops -fprefetch-loop-arrays -fpermissive -fno-exceptions ${PORTABILITY_FLAGS} -pg"
76 export CPPFLAGS=""
77 export CXXCPPFLAGS=""
78 export LDFLAGS="-L${CC_HOME}/lib64 -L${CC_HOME}/lib"
79 export LIBS=""
80 export EXTRA_LIBS=""
81 export PARMACS_MACRO_FILE="threads"
82
```

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ source env.sh
ashen@Stephanie-Lin:~/parsec-benchmark-master$ parsecmgmt -a build -p parsec.dedup -c gcc
[PARSEC] Packages to build: parsec.dedup

[PARSEC] [===== Building package parsec.dedup [1] =====]
[PARSEC] [----- Analyzing package parsec.dedup -----]
[PARSEC] Package parsec.dedup already exists, proceeding.
[PARSEC]
[PARSEC] BIBLIOGRAPHY
[PARSEC]
[PARSEC] [1] Bienia. Benchmarking Modern Multiprocessors. Ph.D. Thesis, 2011.
[PARSEC]
[PARSEC] Done.
```

Dedup build with gcc -O0

```
73# Arguments to use
74export CFLAGS="-DUNIX -O0 -funroll-loops -fprefetch-loop-arrays ${PORTABILITY_FLAGS} -pg"
75export CXXFLAGS="-DUNIX -O0 -funroll-loops -fprefetch-loop-arrays -fpermissive -fno-exceptions ${PORTABILITY_FLAGS} -pg"
76export CPPFLAGS=""
77export CXXCPPFLAGS=""
78export LDFLAGS="-L${CC_HOME}/lib64 -L${CC_HOME}/lib"
79export LIBS=""
80export EXTRA_LIBS=""
81export PARMACS_MACRO_FILE="pthread"
```

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ parsecmgmt -a build -p parsec.dedup -c gcc
[PARSEC] Packages to build: parsec.dedup

[PARSEC] [===== Building package parsec.dedup [1] =====]
[PARSEC] [----- Analyzing package parsec.dedup -----]
[PARSEC] parsec.dedup depends on: ssl zlib
[PARSEC] [----- Analyzing package parsec.ssl -----]
[PARSEC] Package parsec.ssl already exists, proceeding.
[PARSEC] [----- Analyzing package parsec.zlib -----]
[PARSEC] Package parsec.zlib already exists, proceeding.
[PARSEC] [----- Building package parsec.dedup -----]
[PARSEC] Copying source code of package parsec.dedup.
```


Dedup build with icc -O3

```
--  
73 # Arguments to use  
74 export CFLAGS="-gcc-name=${GCC_HOME}/bin/gcc -gcc -I${GCC_HOME}/include/c++/4.2.1 -O3 -funroll-loops -qopt-prefetch"  
75 export CXXFLAGS="-gcc-name=${GCC_HOME}/bin/gcc -gcc -I${GCC_HOME}/include/c++/4.2.1 -O3 -funroll-loops -qopt-prefetch"  
76 export CPPFLAGS=""  
77 export CXXCPPFLAGS=""  
78 export LDFLAGS="-L${CC_HOME}/lib -L${CC_HOME}/lib"  
79 export LIBS=""  
80 export EXTRA_LIBS=""  
--
```

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ parsecmgmt -a build -p parsec.dedup -c icc  
icc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed in a future release. Please  
on to use this compiler. Use '-diag-disable=10441' to disable this message.  
icpc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed in a future release. Please  
ion to use this compiler. Use '-diag-disable=10441' to disable this message.  
[PARSEC] Packages to build: parsec.dedup  
  
[PARSEC] [===== Building package parsec.dedup [1] =====]  
[PARSEC] [----- Analyzing package parsec.dedup -----]  
[PARSEC] parsec.dedup depends on: ssl zlib  
[PARSEC] [----- Analyzing package parsec.ssl -----]  
[PARSEC] Package parsec.ssl already exists, proceeding.  
[PARSEC] [----- Analyzing package parsec.zlib -----]  
[PARSEC] Package parsec.zlib already exists, proceeding.  
[PARSEC] [----- Building package parsec.dedup -----]  
[PARSEC] [===== Building package parsec.dedup [1] =====]
```

Dedup build with icc -O0

```
73 # Arguments to use
74 export CFLAGS="-gcc-name=${GCC_HOME}/bin/gcc -gcc -I${GCC_HOME}/include/c++/4.2.1 -O0 -funroll-loops -qopt-prefetch"
75 export CXXFLAGS="-gcc-name=${GCC_HOME}/bin/gcc -gcc -I${GCC_HOME}/include/c++/4.2.1 -O0 -funroll-loops -qopt-prefetch"
76 export CPPFLAGS=""
77 export CXXCPPFLAGS=""
78 export LDFLAGS="-L${CC_HOME}/lib -L${CC_HOME}/lib"
79 export LIBS=""
80 export EXTRA_LIBS=""
```

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ parsecmgmt -a build -p parsec.dedup -c icc
icc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed in a future release. Please consider moving on to use this compiler. Use '-diag-disable=10441' to disable this message.
icpc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed in a future release. Please consider moving on to use this compiler. Use '-diag-disable=10441' to disable this message.
[PARSEC] Packages to build: parsec.dedup

[PARSEC] [===== Building package parsec.dedup [1] =====]
[PARSEC] [----- Analyzing package parsec.dedup -----]
[PARSEC] parsec.dedup depends on: ssl zlib
[PARSEC] [----- Analyzing package parsec.ssl -----]
[PARSEC] Package parsec.ssl already exists, proceeding.
[PARSEC] [----- Analyzing package parsec.zlib -----]
[PARSEC] Package parsec.zlib already exists, proceeding.
[PARSEC] [----- Building package parsec.dedup -----]
```

Experiment Result

Dedup gcc -O3 run with native test and 4 threads

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ time parsecmgmt -a run -p parsec.dedup -c gcc -i native -n 4
[PARSEC] Benchmarks to run: parsec.dedup

[PARSEC] [===== Running benchmark parsec.dedup [1] =====]
[PARSEC] Setting up run directory.
[PARSEC] Unpacking benchmark input 'native'.
FC-6-x86_64-disc1.iso
[PARSEC] Running ' /home/ashen/parsec-benchmark-master/pkgs/kernels/dedup/inst/amd64-linux.gcc/bin/dedup -c -p -v -t 4 -i FC-6-x86_64-disc1.iso -o output.dat.ddp':
[PARSEC] [----- Beginning of output -----]
PARSEC Benchmark Suite Version 3.0-beta-20150206
Total input size:          671.58 MB
Total output size:         637.28 MB
Effective compression factor: 1.05x

Mean data chunk size:      1.88 KB (stddev: 2023.50 KB)
Amount of duplicate chunks: 54.49%
Data size after deduplication: 658.95 MB (compression factor: 1.02x)
Data size after compression: 630.26 MB (compression factor: 1.05x)
Output overhead:          1.10%
[PARSEC] [----- End of output -----]
[PARSEC]
[PARSEC] BIBLIOGRAPHY
[PARSEC]
[PARSEC] [1] Bienia. Benchmarking Modern Multiprocessors. Ph.D. Thesis, 2011.
[PARSEC]
[PARSEC] Done.

real    0m4.940s
user    0m14.492s
sys     0m3.518s
```

Dedup gcc -O0 run with native test and 4 threads

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ time parsecmgmt -a run -p parsec.dedup -c gcc -i native -n 4
[PARSEC] Benchmarks to run: parsec.dedup

[PARSEC] [===== Running benchmark parsec.dedup [1] =====]
[PARSEC] Setting up run directory.
[PARSEC] Unpacking benchmark input 'native'.
FC-6-x86_64-disc1.iso
[PARSEC] Running ' /home/ashen/parsec-benchmark-master/pkgs/kernels/dedup/inst/amd64-linux.gcc/bin/dedup -c -p -v -t 4 -i FC-6-x86_64-disc1.iso -o output.dat.ddp':
[PARSEC] [----- Beginning of output -----]
PARSEC Benchmark Suite Version 3.0-beta-20150206
Total input size:          671.58 MB
Total output size:         637.28 MB
Effective compression factor: 1.05x

Mean data chunk size:      1.88 KB (stddev: 2023.50 KB)
Amount of duplicate chunks: 54.49%
Data size after deduplication: 658.95 MB (compression factor: 1.02x)
Data size after compression: 630.26 MB (compression factor: 1.05x)
Output overhead:          1.10%
[PARSEC] [----- End of output -----]
[PARSEC]
[PARSEC] BIBLIOGRAPHY
[PARSEC]
[PARSEC] [1] Bienia. Benchmarking Modern Multiprocessors. Ph.D. Thesis, 2011.
[PARSEC]
[PARSEC] Done.

real    0m5.923s
user    0m19.574s
sys     0m5.635s
```


Dedup icc -O3 run with native test and 4 threads

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ time parsecmgmt -a run -p parsec.dedup -c icc -i native -n 4
icc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed from product release in the second half of 2023. The Intel(R) oneAPI DP
on to use this compiler. Use '-diag-disable=10441' to disable this message.
icpc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed from product release in the second half of 2023. The Intel(R) oneAPI DP
ion to use this compiler. Use '-diag-disable=10441' to disable this message.
[PARSEC] Benchmarks to run: parsec.dedup

[PARSEC] [===== Running benchmark parsec.dedup [1] =====]
[PARSEC] Setting up run directory.
[PARSEC] Unpacking benchmark input 'native'.
FC-6-x86_64-disc1.iso
[PARSEC] Running ' /home/ashen/parsec-benchmark-master/pkgs/kernels/dedup/inst/amd64-linux.icc/bin/dedup -c -p -v -t 4 -i FC-6-x86_64-disc1.iso -o output.dat.ddp':
[PARSEC] [----- Beginning of output -----]
PARSEC Benchmark Suite Version 3.0-beta-20150206
Total input size:          671.58 MB
Total output size:         637.28 MB
Effective compression factor: 1.05x

Mean data chunk size:      1.88 KB (stddev: 2023.50 KB)
Amount of duplicate chunks: 54.49%
Data size after deduplication: 658.95 MB (compression factor: 1.02x)
Data size after compression: 630.26 MB (compression factor: 1.05x)
Output overhead:          1.10%
[PARSEC] [----- End of output -----]
[PARSEC]
[PARSEC] BIBLIOGRAPHY
[PARSEC]
[PARSEC] [1] Bienia. Benchmarking Modern Multiprocessors. Ph.D. Thesis, 2011.
[PARSEC]
[PARSEC] Done.

real    0m5.249s
user    0m13.975s
sys     0m3.782s
```

Dedup icc -O0 run with native test and 4 threads

```
ashen@Stephanie-Lin:~/parsec-benchmark-master$ time parsecmgmt -a run -p parsec.dedup -c icc -i native -n 4
icc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed from product release in the second half of 2023. The Intel(R) oneAPI DP
on to use this compiler. Use '-diag-disable=10441' to disable this message.
icpc: remark #10441: The Intel(R) C++ Compiler Classic (ICC) is deprecated and will be removed from product release in the second half of 2023. The Intel(R) oneAPI D
ion to use this compiler. Use '-diag-disable=10441' to disable this message.
[PARSEC] Benchmarks to run: parsec.dedup

[PARSEC] [===== Running benchmark parsec.dedup [1] =====]
[PARSEC] Setting up run directory.
[PARSEC] Unpacking benchmark input 'native'.
FC-6-x86_64-disc1.iso
[PARSEC] Running ' /home/ashen/parsec-benchmark-master/pkgs/kernels/dedup/inst/amd64-linux.icc/bin/dedup -c -p -v -t 4 -i FC-6-x86_64-disc1.iso -o output.dat.ddp':
[PARSEC] [----- Beginning of output -----]
PARSEC Benchmark Suite Version 3.0-beta-20150206
Total input size:          671.58 MB
Total output size:         637.28 MB
Effective compression factor: 1.05x

Mean data chunk size:      1.88 KB (stddev: 2023.50 KB)
Amount of duplicate chunks: 54.49%
Data size after deduplication: 658.95 MB (compression factor: 1.02x)
Data size after compression: 630.26 MB (compression factor: 1.05x)
Output overhead:          1.10%
[PARSEC] [----- End of output -----]
[PARSEC]
[PARSEC] BIBLIOGRAPHY
[PARSEC]
[PARSEC] [1] Bienia. Benchmarking Modern Multiprocessors. Ph.D. Thesis, 2011.
[PARSEC]
[PARSEC] Done.
```

```
real    0m5.344s
user    0m14.020s
sys     0m4.003s
```

Execution Time

- GCC

- -O3

real	0m4.940s
user	0m14.492s
sys	0m3.518s

- -O0

real	0m5.923s
user	0m19.574s
sys	0m5.635s

- ICC

- -O3

real	0m5.249s
user	0m13.975s
sys	0m3.782s

- -O0

real	0m5.344s
user	0m14.020s
sys	0m4.003s

Reference

- <https://github.com/bamos/parsec-benchmark>
- <http://abhishek-sagar.blogspot.com/2012/06/parsec-md5-x8664s41-error-0xd76aa478.html>
- <https://yulistic.gitlab.io/2016/05/parsec-3.0-installation-issues/>